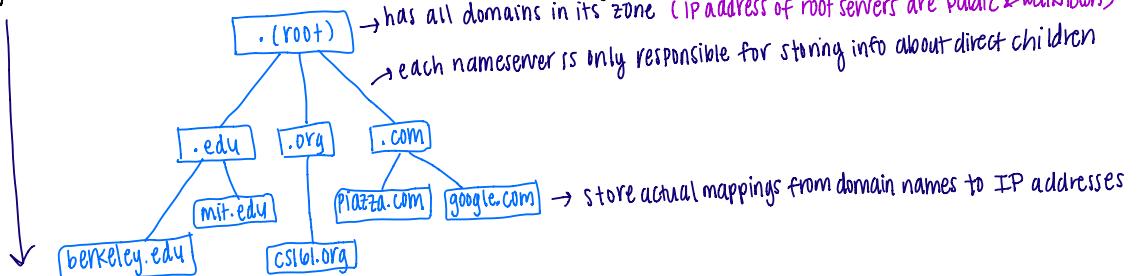


- ↳ shared PS is result of key exchange:  $g^{ab} \bmod p$
- advantage: forward secrecy (even if attacker gets server's private key cannot figure out a or b)
- use PS, R<sub>B</sub>, R<sub>S</sub> to generate set of 4 shared symmetric keys
  - ↳ encryption key C<sub>B</sub> } for client
  - ↳ integrity key I<sub>B</sub> }
  - ↳ encryption key C<sub>S</sub> } for server
  - ↳ integrity key I<sub>S</sub> }
- client & server exchange & verify MACs over all messages sent so far (using I<sub>B</sub> & I<sub>S</sub>) to ensure no one has tampered w/ msgs sent in handshake so far
- now msgs are encrypted & MACed w/ C & I of sender before being sent
  - ↳ msgs have full confidentiality & integrity = TLS has achieved end-to-end security between client & server
- In practice, msgs sent over TLS usually include some counter/timestep so that an attacker cannot record a TLS msg & send it again within the same connection
- biggest advantage & problem of TLS is certificate authorities
  - ↳ good: delegate trust
  - ↳ bad: all CAs need to be trusted to speak for every site

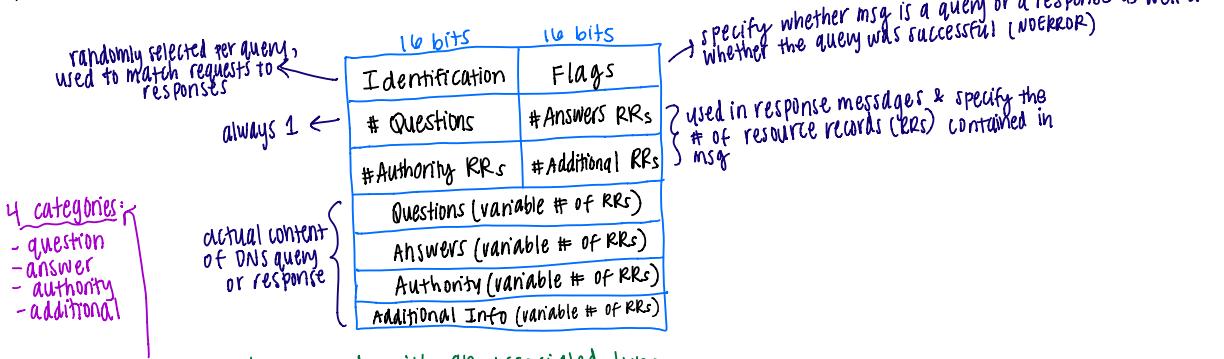
## - DNS (Domain Name System)

- ↳ protocol to translate between human-readable names and IP addresses
- uses a collection of many name servers (servers dedicated to replying to DNS requests)
  - ↳ each one is responsible for specific zone of domains
  - ↳ they also have their own domain names & IP addresses
- name server hierarchy
- when you query a name server, instead of always returning the IP address of the domain you queried, the name server can also direct you to another name



- computer can manually perform DNS lookups but in practice, local computer usually delegates the task of DNS lookups to a DNS Recursive Resolver (provided by ISP)
  - ↳ sends queries, processes responses, & maintains internal cache of records
  - ↳ when performing a lookup, DNS Stub Resolver on computer sends a query to recursive resolver, lets it do all the work, and receives the response
- DNS message format

→ designed to be very lightweight & fast (uses UDP)  
 → packet header



each RR: key / value pair with an associated type

↳ (TTL, value) ↳ actual value data  
 (Name, Class = IN, Type) ↳ time to live record type (A: map domains to IP addresses)  
 ↳ actual key ↳ NS: map zones to domains

## - DNS Security

- DNS is insecure against a malicious name server
  - ↳ to prevent this, resolvers implement **baillwick checking** → a nameserver is only allowed to provide records in its zone
- against on-path attacker, DNS is completely insecure (everything is sent over plaintext)
  - ↳ race condition
- against off-path attacker, **Kaminsky attack**
  - ↳ relies on querying for nonexistent domains
  - legitimate response to nonexistent domain is NXDOMAIN status with no other records (nothing is cached)
    - ↳ attacker can include malicious additional records in fake response for nonexistent domain
      - ↳ race condition (fake vs. real response)
  - to force victim to make thousands of DNS queries for nonexistent domains, trick victim into visiting a website that tries to load lots of nonexistent domains (img src)
- Kaminsky attack allows on-path attackers to race until fake response arrives first & off-path attackers to race until they successfully guess the ID field
  - ↳ UDP source port randomization makes this attack more difficult

## - DNSSEC

- ↳ an extension to regular DNS that provides **integrity & authentication** on all DNS msgs sent
  - every name server generates a public/private key pair and signs every record it sends with its private key
    - ↳ when server receives a DNS request, it sends records, signature on records, and public key to resolver; resolver uses public key to verify signature
- to defend against malicious name servers use **trust anchor** (instead of inherently trusting all name servers)
  - ↳ root server is trust anchor
    - ↳ how can delegate trust from trust anchor to somebody else (by signing their public key)
- new DNSSEC record types (added to DNS design)
  - DNSKEY type record to encode public key
  - RRSIG type record is signature on set of multiple other records in msg (all of same type)
  - DS type record is hash of signer's name & a child's public key (allows trusted server to sign public key of child)
- each DNSSEC name server **generates 2 private/public key pairs**
  - ↳ KSK is only used to sign zone signing key
  - ↳ ZSK is used to sign everything else
    - ↳ "parent" endorses by signing it  
"child"
- for DNS & DNSSEC, start by first querying for root (hardcoded IP address)
- **DNS is meant to be fast, DNSSEC is slow**
  - ↳ DNSSEC signs records offline & saves signatures w/ records in server
    - ↳ to handle nonexistent domains (w/out triggering DoS), name servers precompute signatures on ranges of nonexistent domains (use hashes instead of actual names for better security)

## • Web security

- every resource on the web is identified by a URL

↳ `http://www.example.com/index.html`

protocol → tells browser which page on web server to request  
domain name → tells browser which web server to contact to retrieve the resource  
↳ tells browser how to retrieve the resource (HTTP or HTTPS)  
→ secure version of HTTP using TLS

(sometimes include port # to distinguish between diff applications running on same web server)

### - HTTP (Hypertext Transfer Protocol)

- language clients use to communicate with servers in order to fetch resources & issue other requests (can only return plain text files)

#### - Request Response Model

↳ clients (such as browsers) must actively start a connection to the server and issue a request which the server then responds to

#### - Structure of a Request

Line 1: method of request (GET / POST), path of request (/), protocol version (HTTP/1.1)

→ every other line is a request header for GET; POST can include a body as well

#### - GET vs. POST

- GET: intended for "getting" information from the server and generally do not change anything on the server's end

- POST: intended for sending information to the server that somehow modifies its internal state (e.g. adding a comment in a forum or changing your password)

### - Elements of a webpage (3 different languages)

- HTML: lets us create structured documents w/ paragraphs, links, fillable forms, and embedded images, etc.

- frames pose a security risk since the outer page is now including an inner page that may be from diff (malicious) source

↳ to protect against this, browsers enforce frame isolation (outer page cannot change contents of inner page & inner page cannot change contents of outer page)

- CSS: lets us modify the appearance of an HTML page by using diff fonts, colors, spacing, etc.

- Javascript: lets us have dynamic features such as interactive buttons

↳ runs in browser and can arbitrarily modify any HTML or CSS on a webpage

↳ modern web browsers typically run Javascript in a sandbox so that any code from a webpage cannot access sensitive data on your computer

### - Same Origin Policy

- browsing multiple webpages poses a security risk (cuz one may be malicious)

↳ defend against this by enforcing same origin policy (isolates every webpage in your browser except for when two webpages have the same origin)

#### - origins (determined by protocol, domain, and port)

↳ to check if two webpages have the same origin, the same-origin policy performs string matching on protocol, domain, and port

↳ same origin if all 3 match (if port is not specified, defaults to 80) ↳ default ports are the ones exception to string matching

#### - Exceptions (to origin of webpage is defined by its URL rule)

- Javascript runs with the origin of the page that loads it

i.e.) `<script src="http://abc.com"></script>` on `http://cs.org`

↳ script has origin of `http://cs.org`

- Images have origin of page that loads it

- Frames have origin of URL where the frame is retrieved from, not the origin of the website that loads it

i.e.) `<iframe src="http://google.com"></iframe>` on `http://cs161.org`

↳ frame has origin of `http://google.com`

### - SQL Injection

- If web server uses user input as part of the code it runs and the input is not properly checked, an attacker could create a special input that causes unintended code to run on the server

- trick to force SQL query to always return something  
↳ logic that always returns true (ex. 'OR 1=1')

- Defenses
  - escape inputs
    - ↳ tell SQL to treat every character that could be used in an attack as part of the string (not actual SQL syntax)
    - ↳ but building a good escapert can be tricky
  - parameterized SQL/prepared statements
    - ↳ compiles query first and then plugs in user input after the query has already been interpreted by SQL parser
    - ↳ no way for attacker input to be interpreted as SQL code
    - ↳ best defense against SQL injection
- XSS (cross-site scripting)
  - attacker injects malicious javascript onto a webpage
    - ↳ when victim user loads webpage, user's browser will run malicious javascript
    - ↳ subvert same-origin policy
  - stored XSS: attacker finds way to persistently store malicious javascript on web server; when victim loads the webpage, web server will load malicious JS & display it
    - ↳ make fb post w/ content: <script> alert("lol") </script>
  - reflected XSS: attacker finds a vulnerable webpage where the server receives user input in an HTTP request and displays user input in response
    - ↳ google 'search?<script>...</script>' in URL
- Defenses
  - Sanitize Input
    - ↳ check for malicious input that might cause JavaScript to run like '<script>'
    - ↳ hard to write good detector that catches all attacks
    - ↳ replace potentially dangerous characters with their HTML encoding (< -> &lt;)
  - CSP (content security policy)
    - ↳ specifies a list of allowed domains where scripts can be loaded from
    - ↳ defined by a web server and enforced by a browser
      - ↳ In HTTP response, the server attaches a CSP header and the browser checks any scripts against the header
- Cookies & Session Management
  - HTTP is a stateless protocol, each request and response is isolated from all other requests and responses
  - browsers and servers store HTTP cookies to maintain some form of state
    - ↳ cookie → piece of data stored in browser memory state that should persist across multiple requests and responses
    - ↳ when making a change across states, server sends response with set-cookie header (tells browser to store a new cookie)
      - ↳ in future requests, browser will automatically attach relevant cookies to a request & send it to web server so server can customize response
  - every cookie is a name, value pair (ie. darkmode, True)
  - only some cookies in each request (browser uses cookie attributes to determine which to send)
    - ↳ Domain & Path attributes tell browser which URLs to send cookie to
      - ↳ browser sends a cookie to a given URL if the cookie's domain attribute is a domain suffix of the URL domain & the cookie's path attribute is a prefix of the URL path
        - ↳ URL domain should end in cookie domain & URL path begins w/ cookie path
    - ↳ when a server sets a cookie, the cookie's domain must be a URL suffix of the server's URL
      - ↳ for the cookie to be set, the server's URL must end in the cookie's domain attribute (else, browser will reject cookie)
      - ↳ Exception: cookies cannot have domains set to a top-level domain, such as .edu or .com since there are too broad and pose a security risk
        - (cookie policy allows server to set Path attribute without any restrictions)

- Secure attribute tells browser to only send cookie over secure HTTPS connection
- HttpOnly attribute prevents JS from accessing & modifying the cookie
- expires attribute tells browser when to stop remembering the cookie
- **Session Management**
  - ↳ Cookies are often used to keep users logged in to a website over many requests and responses
  - ↳ When user sends a login request w/ valid username & password, server will generate new session token and send it to user as a cookie
  - ↳ In future requests, the browser will attach session token cookie & send it to the server (server maintains a mapping of session tokens to users so when it receives a request w/ ID, it can look up corresponding user & customize response accordingly)
- Secure session tokens should be random and unpredictable (HttpOnly & secure protect them from XSS vulnerabilities and network attackers)
- **Session token v. cookie**
  - ↳ Session token: values that browser sends to the server to associate the request with a logged-in user
  - ↳ Cookie: how browser stores and sends session tokens to the server (can also save other state)
  - ④ TLDR: session tokens are a special type of cookie that keeps users logged in over many requests & responses
- **CSRF (Cross-site Request Forgery)**
  - ↳ Attacker forces victim to make an unintended request (victim's browser will automatically attach the session token cookie to the unintended request & server will accept request as coming from the victim)
  - Most common over HTTP POST requests (embedded in HTML in malicious link)
    - ↳ Common example that generates HTTP POST request is HTML forms (<form name=>...</form>)
  - Defenses (use both together → defense in depth)
    - **CSRF token** (included on webpages)
      - ↳ When legitimate user loads a webpage from the server with a form, the server will randomly generate a CSRF token & include it as an extra field in the form
      - ↳ When user submits form, the form will include CSRF token and the server will check that the CSRF token is valid (if it's missing or invalid, server rejects request)
    - Need to generate new CSRF token everytime user requests a form
      - ↳ Should be random & unpredictable
      - ↳ Server needs to maintain mapping of CSRF tokens to session tokens
  - **Referer Validation**
    - HTTP request includes Referer header which indicates which URL request was made from
    - Server can check Referer header on each request & reject any requests that have untrusted or suspicious Referer headers
    - Good defense if included on every request (but Referer header could be blank)